

Assignment 1

May 15, 2019

You are currently looking at **version 1.3** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

1 Assignment 1 - Introduction to Machine Learning

For this assignment, you will be using the Breast Cancer Wisconsin (Diagnostic) Database to create a classifier that can help diagnose patients. First, read through the description of the dataset (below).

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import load_breast_cancer

cancer = load_breast_cancer()

#print(cancer.DESCR) # Print the data set description
```

The object returned by `load_breast_cancer()` is a scikit-learn Bunch object, which is similar to a dictionary.

```
In [2]: cancer.keys()

Out[2]: dict_keys(['data', 'target', 'target_names', 'DESCR', 'feature_names'])
```

1.0.1 Question 0 (Example)

How many features does the breast cancer dataset have?

This function should return an integer.

```
In [3]: # You should write your whole answer within the function provided. The auto
# this function and compare the return value against the correct solution v
def answer_zero():
    # This function returns the number of features of the breast cancer dat
    # The assignment question description will tell you the general format
```

```

    return len(cancer['feature_names'])

# You can examine what your function returns by calling it in the cell. If
# about the assignment formats, check out the discussion forums for any FAQ
answer_zero()

```

Out [3]: 30

1.0.2 Question 1

Scikit-learn works with lists, numpy arrays, scipy-sparse matrices, and pandas DataFrames, so converting the dataset to a DataFrame is not necessary for training this model. Using a DataFrame does however help make many things easier such as munging data, so let's practice creating a classifier with a pandas DataFrame.

Convert the sklearn.dataset cancer to a DataFrame.

This function should return a (569, 31) DataFrame with columns =

```

['mean radius', 'mean texture', 'mean perimeter', 'mean area',
 'mean smoothness', 'mean compactness', 'mean concavity',
 'mean concave points', 'mean symmetry', 'mean fractal dimension',
 'radius error', 'texture error', 'perimeter error', 'area error',
 'smoothness error', 'compactness error', 'concavity error',
 'concave points error', 'symmetry error', 'fractal dimension error',
 'worst radius', 'worst texture', 'worst perimeter', 'worst area',
 'worst smoothness', 'worst compactness', 'worst concavity',
 'worst concave points', 'worst symmetry', 'worst fractal dimension',
 'target']

```

and index =

```

RangeIndex(start=0, stop=569, step=1)

```

In [4]: `def answer_one():`

```

    result = pd.DataFrame(data=np.c_[cancer['data'], cancer['target']],
                          columns=np.append(cancer['feature_names'], ['target']))
    return result

```

```

answer_one()

```

```

Out [4]:
   mean radius  mean texture  mean perimeter  mean area  mean smoothness
0         17.990         10.38         122.80      1001.0         0.11840
1         20.570         17.77         132.90      1326.0         0.08474
2         19.690         21.25         130.00      1203.0         0.10960
3         11.420         20.38          77.58       386.1         0.14250
4         20.290         14.34         135.10      1297.0         0.10030
5         12.450         15.70          82.57       477.1         0.12780
6         18.250         19.98         119.60      1040.0         0.09463

```

7	13.710	20.83	90.20	577.9	0.11890
8	13.000	21.82	87.50	519.8	0.12730
9	12.460	24.04	83.97	475.9	0.11860
10	16.020	23.24	102.70	797.8	0.08206
11	15.780	17.89	103.60	781.0	0.09710
12	19.170	24.80	132.40	1123.0	0.09740
13	15.850	23.95	103.70	782.7	0.08401
14	13.730	22.61	93.60	578.3	0.11310
15	14.540	27.54	96.73	658.8	0.11390
16	14.680	20.13	94.74	684.5	0.09867
17	16.130	20.68	108.10	798.8	0.11700
18	19.810	22.15	130.00	1260.0	0.09831
19	13.540	14.36	87.46	566.3	0.09779
20	13.080	15.71	85.63	520.0	0.10750
21	9.504	12.44	60.34	273.9	0.10240
22	15.340	14.26	102.50	704.4	0.10730
23	21.160	23.04	137.20	1404.0	0.09428
24	16.650	21.38	110.00	904.6	0.11210
25	17.140	16.40	116.00	912.7	0.11860
26	14.580	21.53	97.41	644.8	0.10540
27	18.610	20.25	122.10	1094.0	0.09440
28	15.300	25.27	102.40	732.4	0.10820
29	17.570	15.05	115.00	955.1	0.09847
..
539	7.691	25.44	48.34	170.4	0.08668
540	11.540	14.44	74.65	402.9	0.09984
541	14.470	24.99	95.81	656.4	0.08837
542	14.740	25.42	94.70	668.6	0.08275
543	13.210	28.06	84.88	538.4	0.08671
544	13.870	20.70	89.77	584.8	0.09578
545	13.620	23.23	87.19	573.2	0.09246
546	10.320	16.35	65.31	324.9	0.09434
547	10.260	16.58	65.85	320.8	0.08877
548	9.683	19.34	61.05	285.7	0.08491
549	10.820	24.21	68.89	361.6	0.08192
550	10.860	21.48	68.51	360.5	0.07431
551	11.130	22.44	71.49	378.4	0.09566
552	12.770	29.43	81.35	507.9	0.08276
553	9.333	21.94	59.01	264.0	0.09240
554	12.880	28.92	82.50	514.3	0.08123
555	10.290	27.61	65.67	321.4	0.09030
556	10.160	19.59	64.73	311.7	0.10030
557	9.423	27.88	59.26	271.3	0.08123
558	14.590	22.68	96.39	657.1	0.08473
559	11.510	23.93	74.52	403.5	0.09261
560	14.050	27.15	91.38	600.4	0.09929
561	11.200	29.37	70.67	386.0	0.07449
562	15.220	30.62	103.40	716.9	0.10480

563	20.920	25.09	143.00	1347.0	0.10990
564	21.560	22.39	142.00	1479.0	0.11100
565	20.130	28.25	131.20	1261.0	0.09780
566	16.600	28.08	108.30	858.1	0.08455
567	20.600	29.33	140.10	1265.0	0.11780
568	7.760	24.54	47.92	181.0	0.05263

	mean compactness	mean concavity	mean concave points	mean symmetry
0	0.27760	0.300100	0.147100	0.2419
1	0.07864	0.086900	0.070170	0.1812
2	0.15990	0.197400	0.127900	0.2069
3	0.28390	0.241400	0.105200	0.2597
4	0.13280	0.198000	0.104300	0.1809
5	0.17000	0.157800	0.080890	0.2087
6	0.10900	0.112700	0.074000	0.1794
7	0.16450	0.093660	0.059850	0.2196
8	0.19320	0.185900	0.093530	0.2350
9	0.23960	0.227300	0.085430	0.2030
10	0.06669	0.032990	0.033230	0.1528
11	0.12920	0.099540	0.066060	0.1842
12	0.24580	0.206500	0.111800	0.2397
13	0.10020	0.099380	0.053640	0.1847
14	0.22930	0.212800	0.080250	0.2069
15	0.15950	0.163900	0.073640	0.2303
16	0.07200	0.073950	0.052590	0.1586
17	0.20220	0.172200	0.102800	0.2164
18	0.10270	0.147900	0.094980	0.1582
19	0.08129	0.066640	0.047810	0.1885
20	0.12700	0.045680	0.031100	0.1967
21	0.06492	0.029560	0.020760	0.1815
22	0.21350	0.207700	0.097560	0.2521
23	0.10220	0.109700	0.086320	0.1769
24	0.14570	0.152500	0.091700	0.1995
25	0.22760	0.222900	0.140100	0.3040
26	0.18680	0.142500	0.087830	0.2252
27	0.10660	0.149000	0.077310	0.1697
28	0.16970	0.168300	0.087510	0.1926
29	0.11570	0.098750	0.079530	0.1739
..
539	0.11990	0.092520	0.013640	0.2037
540	0.11200	0.067370	0.025940	0.1818
541	0.12300	0.100900	0.038900	0.1872
542	0.07214	0.041050	0.030270	0.1840
543	0.06877	0.029870	0.032750	0.1628
544	0.10180	0.036880	0.023690	0.1620
545	0.06747	0.029740	0.024430	0.1664
546	0.04994	0.010120	0.005495	0.1885
547	0.08066	0.043580	0.024380	0.1669

548	0.05030	0.023370	0.009615	0.1580
549	0.06602	0.015480	0.008160	0.1976
550	0.04227	0.000000	0.000000	0.1661
551	0.08194	0.048240	0.022570	0.2030
552	0.04234	0.019970	0.014990	0.1539
553	0.05605	0.039960	0.012820	0.1692
554	0.05824	0.061950	0.023430	0.1566
555	0.07658	0.059990	0.027380	0.1593
556	0.07504	0.005025	0.011160	0.1791
557	0.04971	0.000000	0.000000	0.1742
558	0.13300	0.102900	0.037360	0.1454
559	0.10210	0.111200	0.041050	0.1388
560	0.11260	0.044620	0.043040	0.1537
561	0.03558	0.000000	0.000000	0.1060
562	0.20870	0.255000	0.094290	0.2128
563	0.22360	0.317400	0.147400	0.2149
564	0.11590	0.243900	0.138900	0.1726
565	0.10340	0.144000	0.097910	0.1752
566	0.10230	0.092510	0.053020	0.1590
567	0.27700	0.351400	0.152000	0.2397
568	0.04362	0.000000	0.000000	0.1587

	mean fractal dimension	...	worst texture	worst perimeter	\
0	0.07871	...	17.33	184.60	
1	0.05667	...	23.41	158.80	
2	0.05999	...	25.53	152.50	
3	0.09744	...	26.50	98.87	
4	0.05883	...	16.67	152.20	
5	0.07613	...	23.75	103.40	
6	0.05742	...	27.66	153.20	
7	0.07451	...	28.14	110.60	
8	0.07389	...	30.73	106.20	
9	0.08243	...	40.68	97.65	
10	0.05697	...	33.88	123.80	
11	0.06082	...	27.28	136.50	
12	0.07800	...	29.94	151.70	
13	0.05338	...	27.66	112.00	
14	0.07682	...	32.01	108.80	
15	0.07077	...	37.13	124.10	
16	0.05922	...	30.88	123.40	
17	0.07356	...	31.48	136.80	
18	0.05395	...	30.88	186.80	
19	0.05766	...	19.26	99.70	
20	0.06811	...	20.49	96.09	
21	0.06905	...	15.66	65.13	
22	0.07032	...	19.08	125.10	
23	0.05278	...	35.59	188.00	
24	0.06330	...	31.56	177.00	

25	0.07413	...	21.40	152.40
26	0.06924	...	33.21	122.40
27	0.05699	...	27.26	139.90
28	0.06540	...	36.71	149.30
29	0.06149	...	19.52	134.90
..
539	0.07751	...	31.89	54.49
540	0.06782	...	19.68	78.78
541	0.06341	...	31.73	113.50
542	0.05680	...	32.29	107.40
543	0.05781	...	37.17	92.48
544	0.06688	...	24.75	99.17
545	0.05801	...	29.09	97.58
546	0.06201	...	21.77	71.12
547	0.06714	...	22.04	71.08
548	0.06235	...	25.59	69.10
549	0.06328	...	31.45	83.90
550	0.05948	...	24.77	74.08
551	0.06552	...	28.26	77.80
552	0.05637	...	36.00	88.10
553	0.06576	...	25.05	62.86
554	0.05708	...	35.74	88.84
555	0.06127	...	34.91	69.57
556	0.06331	...	22.88	67.88
557	0.06059	...	34.24	66.50
558	0.06147	...	27.27	105.90
559	0.06570	...	37.16	82.28
560	0.06171	...	33.17	100.20
561	0.05502	...	38.30	75.19
562	0.07152	...	42.79	128.70
563	0.06879	...	29.41	179.10
564	0.05623	...	26.40	166.10
565	0.05533	...	38.25	155.00
566	0.05648	...	34.12	126.70
567	0.07016	...	39.42	184.60
568	0.05884	...	30.37	59.16

	worst area	worst smoothness	worst compactness	worst concavity	\
0	2019.0	0.16220	0.66560	0.71190	
1	1956.0	0.12380	0.18660	0.24160	
2	1709.0	0.14440	0.42450	0.45040	
3	567.7	0.20980	0.86630	0.68690	
4	1575.0	0.13740	0.20500	0.40000	
5	741.6	0.17910	0.52490	0.53550	
6	1606.0	0.14420	0.25760	0.37840	
7	897.0	0.16540	0.36820	0.26780	
8	739.3	0.17030	0.54010	0.53900	
9	711.4	0.18530	1.05800	1.10500	

10	1150.0	0.11810	0.15510	0.14590
11	1299.0	0.13960	0.56090	0.39650
12	1332.0	0.10370	0.39030	0.36390
13	876.5	0.11310	0.19240	0.23220
14	697.7	0.16510	0.77250	0.69430
15	943.2	0.16780	0.65770	0.70260
16	1138.0	0.14640	0.18710	0.29140
17	1315.0	0.17890	0.42330	0.47840
18	2398.0	0.15120	0.31500	0.53720
19	711.2	0.14400	0.17730	0.23900
20	630.5	0.13120	0.27760	0.18900
21	314.9	0.13240	0.11480	0.08867
22	980.9	0.13900	0.59540	0.63050
23	2615.0	0.14010	0.26000	0.31550
24	2215.0	0.18050	0.35780	0.46950
25	1461.0	0.15450	0.39490	0.38530
26	896.9	0.15250	0.66430	0.55390
27	1403.0	0.13380	0.21170	0.34460
28	1269.0	0.16410	0.61100	0.63350
29	1227.0	0.12550	0.28120	0.24890
..
539	223.6	0.15960	0.30640	0.33930
540	457.8	0.13450	0.21180	0.17970
541	808.9	0.13400	0.42020	0.40400
542	826.4	0.10600	0.13760	0.16110
543	629.6	0.10720	0.13810	0.10620
544	688.6	0.12640	0.20370	0.13770
545	729.8	0.12160	0.15170	0.10490
546	384.9	0.12850	0.08842	0.04384
547	357.4	0.14610	0.22460	0.17830
548	364.2	0.11990	0.09546	0.09350
549	505.6	0.12040	0.16330	0.06194
550	412.3	0.10010	0.07348	0.00000
551	436.6	0.10870	0.17820	0.15640
552	594.7	0.12340	0.10640	0.08653
553	295.8	0.11030	0.08298	0.07993
554	595.7	0.12270	0.16200	0.24390
555	357.6	0.13840	0.17100	0.20000
556	347.3	0.12650	0.12000	0.01005
557	330.6	0.10730	0.07158	0.00000
558	733.5	0.10260	0.31710	0.36620
559	474.2	0.12980	0.25170	0.36300
560	706.7	0.12410	0.22640	0.13260
561	439.6	0.09267	0.05494	0.00000
562	915.0	0.14170	0.79170	1.17000
563	1819.0	0.14070	0.41860	0.65990
564	2027.0	0.14100	0.21130	0.41070
565	1731.0	0.11660	0.19220	0.32150

566	1124.0	0.11390	0.30940	0.34030
567	1821.0	0.16500	0.86810	0.93870
568	268.6	0.08996	0.06444	0.00000

	worst concave points	worst symmetry	worst fractal dimension	target
0	0.26540	0.4601	0.11890	0.0
1	0.18600	0.2750	0.08902	0.0
2	0.24300	0.3613	0.08758	0.0
3	0.25750	0.6638	0.17300	0.0
4	0.16250	0.2364	0.07678	0.0
5	0.17410	0.3985	0.12440	0.0
6	0.19320	0.3063	0.08368	0.0
7	0.15560	0.3196	0.11510	0.0
8	0.20600	0.4378	0.10720	0.0
9	0.22100	0.4366	0.20750	0.0
10	0.09975	0.2948	0.08452	0.0
11	0.18100	0.3792	0.10480	0.0
12	0.17670	0.3176	0.10230	0.0
13	0.11190	0.2809	0.06287	0.0
14	0.22080	0.3596	0.14310	0.0
15	0.17120	0.4218	0.13410	0.0
16	0.16090	0.3029	0.08216	0.0
17	0.20730	0.3706	0.11420	0.0
18	0.23880	0.2768	0.07615	0.0
19	0.12880	0.2977	0.07259	1.0
20	0.07283	0.3184	0.08183	1.0
21	0.06227	0.2450	0.07773	1.0
22	0.23930	0.4667	0.09946	0.0
23	0.20090	0.2822	0.07526	0.0
24	0.20950	0.3613	0.09564	0.0
25	0.25500	0.4066	0.10590	0.0
26	0.27010	0.4264	0.12750	0.0
27	0.14900	0.2341	0.07421	0.0
28	0.20240	0.4027	0.09876	0.0
29	0.14560	0.2756	0.07919	0.0
..
539	0.05000	0.2790	0.10660	1.0
540	0.06918	0.2329	0.08134	1.0
541	0.12050	0.3187	0.10230	1.0
542	0.10950	0.2722	0.06956	1.0
543	0.07958	0.2473	0.06443	1.0
544	0.06845	0.2249	0.08492	1.0
545	0.07174	0.2642	0.06953	1.0
546	0.02381	0.2681	0.07399	1.0
547	0.08333	0.2691	0.09479	1.0
548	0.03846	0.2552	0.07920	1.0
549	0.03264	0.3059	0.07626	1.0
550	0.00000	0.2458	0.06592	1.0

551	0.06413	0.3169	0.08032	1.0
552	0.06498	0.2407	0.06484	1.0
553	0.02564	0.2435	0.07393	1.0
554	0.06493	0.2372	0.07242	1.0
555	0.09127	0.2226	0.08283	1.0
556	0.02232	0.2262	0.06742	1.0
557	0.00000	0.2475	0.06969	1.0
558	0.11050	0.2258	0.08004	1.0
559	0.09653	0.2112	0.08732	1.0
560	0.10480	0.2250	0.08321	1.0
561	0.00000	0.1566	0.05905	1.0
562	0.23560	0.4089	0.14090	0.0
563	0.25420	0.2929	0.09873	0.0
564	0.22160	0.2060	0.07115	0.0
565	0.16280	0.2572	0.06637	0.0
566	0.14180	0.2218	0.07820	0.0
567	0.26500	0.4087	0.12400	0.0
568	0.00000	0.2871	0.07039	1.0

[569 rows x 31 columns]

1.0.3 Question 2

What is the class distribution? (i.e. how many instances of malignant (encoded 0) and how many benign (encoded 1)?)

This function should return a Series named target of length 2 with integer values and index = ['malignant', 'benign']

```
In [5]: def answer_two():
        cancerdf = answer_one()

        # Your code here
        cancerdf = answer_one()

        malignant = int(cancerdf[cancerdf['target'] == 0]['target'].count())
        benign = int(cancerdf[cancerdf['target'] == 1]['target'].count())

        result = pd.Series({'malignant': malignant, 'benign': benign})

        return result
```

```
answer_two()
```

```
Out[5]: benign      357
        malignant    212
        dtype: int64
```

1.0.4 Question 3

Split the DataFrame into X (the data) and y (the labels).

*This function should return a tuple of length 2: (X, y), where * X, a pandas DataFrame, has shape (569, 30) * y, a pandas Series, has shape (569,).*

```
In [7]: def answer_three():
        cancerdf = answer_one()

        # Your code here

        X = cancerdf.drop('target', axis=1)
        y = cancerdf['target']
        return X, y
```

1.0.5 Question 4

Using train_test_split, split X and y into training and test sets (X_train, X_test, y_train, and y_test).

Set the random number generator state to 0 using random_state=0 to make sure your results match the autograder!

*This function should return a tuple of length 4: (X_train, X_test, y_train, y_test), where * X_train has shape (426, 30) * X_test has shape (143, 30) * y_train has shape (426,) * y_test has shape (143,)*

```
In [8]: from sklearn.model_selection import train_test_split

        def answer_four():
            X, y = answer_three()

            # Your code here
            X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=426)
            return X_train, X_test, y_train, y_test
```

1.0.6 Question 5

Using KNeighborsClassifier, fit a k-nearest neighbors (knn) classifier with X_train, y_train and using one nearest neighbor (n_neighbors = 1).

This function should return a sklearn.neighbors.classification.KNeighborsClassifier.

```
In [9]: from sklearn.neighbors import KNeighborsClassifier

        def answer_five():
            X_train, X_test, y_train, y_test = answer_four()

            # Your code here
            neigh = KNeighborsClassifier(n_neighbors=1)
            result = neigh.fit(X_train, y_train)
            return result
```

1.0.7 Question 6

Using your knn classifier, predict the class label using the mean value for each feature.

Hint: You can use `cancerdf.mean()[:-1].values.reshape(1, -1)` which gets the mean value for each feature, ignores the target column, and reshapes the data from 1 dimension to 2 (necessary for the `predict` method of `KNeighborsClassifier`).

This function should return a numpy array either `array([0.])` or `array([1.])`

```
In [11]: def answer_six():
          cancerdf = answer_one()
          means = cancerdf.mean()[:-1].values.reshape(1, -1)

          # Your code here
          result = answer_five().predict(means)

          return result
```

1.0.8 Question 7

Using your knn classifier, predict the class labels for the test set `X_test`.

This function should return a numpy array with shape `(143,)` and values either `0.0` or `1.0`.

```
In [13]: def answer_seven():
          X_train, X_test, y_train, y_test = answer_four()
          knn = answer_five()

          # Your code here
          result = knn.predict(X_test)

          return result
```

1.0.9 Question 8

Find the score (mean accuracy) of your knn classifier using `X_test` and `y_test`.

This function should return a float between 0 and 1

```
In [14]: def answer_eight():
          X_train, X_test, y_train, y_test = answer_four()
          knn = answer_five()

          # Your code here
          result = knn.score(X_test, y_test)

          return result
```

1.0.10 Optional plot

Try using the plotting function below to visualize the different prediction scores between training and test sets, as well as malignant and benign cells.

```

In [15]: def accuracy_plot():
import matplotlib.pyplot as plt

%matplotlib notebook

X_train, X_test, y_train, y_test = answer_four()

# Find the training and testing accuracies by target value (i.e. malignant/benign)
mal_train_X = X_train[y_train==0]
mal_train_y = y_train[y_train==0]
ben_train_X = X_train[y_train==1]
ben_train_y = y_train[y_train==1]

mal_test_X = X_test[y_test==0]
mal_test_y = y_test[y_test==0]
ben_test_X = X_test[y_test==1]
ben_test_y = y_test[y_test==1]

knn = answer_five()

scores = [knn.score(mal_train_X, mal_train_y), knn.score(ben_train_X,
                                                         knn.score(mal_test_X, mal_test_y), knn.score(ben_test_X, ben

plt.figure()

# Plot the scores as a bar chart
bars = plt.bar(np.arange(4), scores, color=['#4c72b0', '#4c72b0', '#55a868', '#55a868'])

# directly label the score onto the bars
for bar in bars:
    height = bar.get_height()
    plt.gca().text(bar.get_x() + bar.get_width()/2, height*.90, '{0:.1f}'.format(score))
    ha='center', color='w', fontsize=11)

# remove all the ticks (both axes), and tick labels on the Y axis
plt.tick_params(top='off', bottom='off', left='off', right='off', labelbottom='off')

# remove the frame of the chart
for spine in plt.gca().spines.values():
    spine.set_visible(False)

plt.xticks([0,1,2,3], ['Malignant\nTraining', 'Benign\nTraining', 'Malignant\nTest', 'Benign\nTest'])
plt.title('Training and Test Accuracies for Malignant and Benign Cells')

```

Uncomment the plotting function to see the visualization.

Comment out the plotting function when submitting your notebook for grading.

```

In [17]: #accuracy_plot()

```

```
In [ ]:
```